

Case Study: NoC Latency Estimation


Marcel Mettler

Technische Universität München

Department of Electrical and Computer Engineering

Chair of Electronic Design Automation

25. January 2019




Outline

- Background
 - Introduction to Networks-on-Chip
 - Optimization of Network-on-Chip Topologies
- Application of Neural Networks
- Introduction to the ML framework

Network-on-Chip

On-Chip interconnect inspired by a computer network

- **Processing Element**
Computation unit e.g. CPU, Video Decoder
- **Network Interface**
Translates communication protocol between network and PE
- **Router**
Forward packets through the network according to a routing policy
- **Link**
Point-to-point connections between routers




Network-on-Chip

Benefits

- Scalability
 - Overall bandwidth increases with the network size
 - Especially suitable for a high number of PEs
- Short point-to-point connections
 - Enables high clock frequencies
- Segmentation
 - Increases reusability

Drawbacks

- Latency
 - Multiple hops and multiple cycles per hop
- Chip Area



Network-on-Chip Performance Metrics

- **Aggregated Bandwidth**

The product of the bandwidths per link bw_l and the number of links n_l

$$BW_{agg} = n_l * bw_l$$

- **Throughput**

Usable share of the aggregated sending bandwidth

- **Average Latency**


The average difference between the reception time t_{rx} and the transmit time t_{tx} of all packets $p \in P$

$$l_{avg} = \frac{1}{|P|} \sum_{p \in P} t_{rx}(p) - t_{tx}(p)$$

- **Chip Area**


- **Power Consumption**

Regular Network-on-Chip Topologies



- **High Connectivity**



- E.g. Mesh, Torus
- High throughput and low latency
- High power and area consumption



- **Low Connectivity**

- E.g. Ring
- Low power and area consumption
- Low throughput and high latencies

General-purpose vs. Application-specific NoCs



General-purpose NoCs

- Suited for general-purpose MPSoCs
- Most PEs are not fixed in functionality
 - Traffic requirements are unknown at design time
 - Regular network topologies beneficial

Application-specific NoCs

- Suited for application-specific MPSoCs
- Most PEs are fixed in their functionality
 - Traffic requirements are known at design time
 - **Topology must be optimized w.r.t. application**

Outline

- Background
 - Introduction to Networks-on-Chip
 - Optimization of Network-on-Chip Topologies
- Application of Neural Networks
- Introduction to the ML framework

Optimization Networks-on-Chip Topologies

Design task: Find the best-suited NoC topology for the application at hand

- The quality of a NoC topology x for an application t

$$Q(x, t) = -1 * [w_l l_{avg}(x, t) + w_P P(x) + w_A A(x)], \text{ where}$$

$l_{avg}(x, t) :=$ average latency

$P(x) :=$ power dissipation

$A(x) :=$ chip area

$x :=$ NoC topology


$t :=$ application

→ Evaluation toolchain required to evaluate $l_{avg}(x, t)$, $P(x)$ and $A(x)$

NoC Evaluation Toolchain

Application t

- Described by Core Communication Graph



Write Video Data from
DDR Memory to SD Card

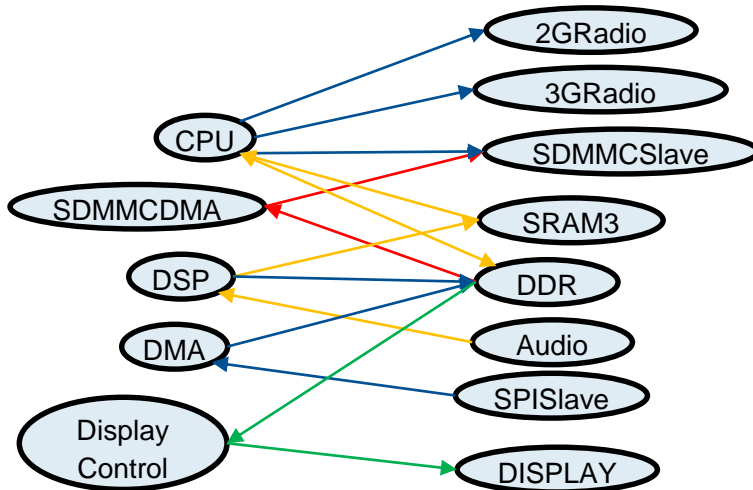
Record Audio

Show Video

NoC Evaluation Toolchain

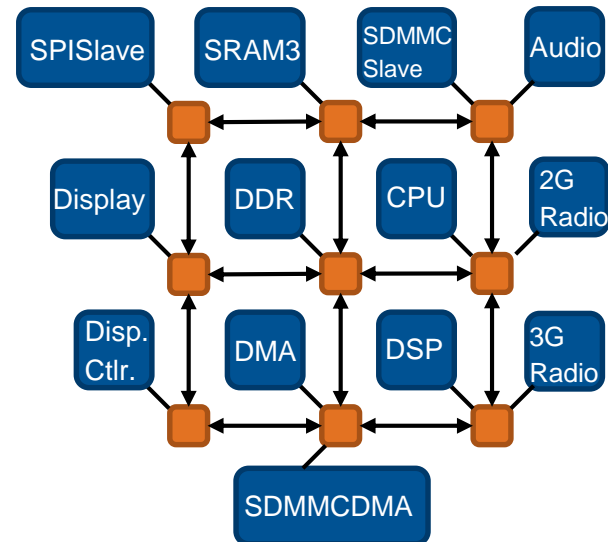
Application t

- Described by Core Communication Graph



NoC Topology x

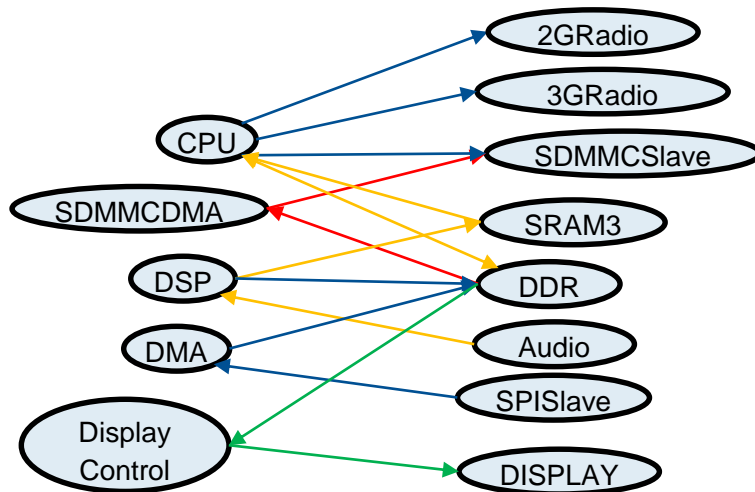
- Described the topology graph of the network



NoC Evaluation Toolchain


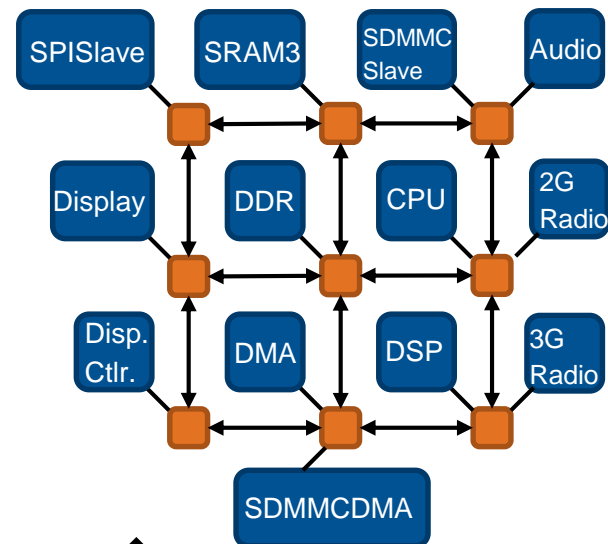
Application t

- Described by Core Communication Graph



NoC Topology x

- Described the topology graph of the network



NoC Evaluation Toolchain

ORION[2]:

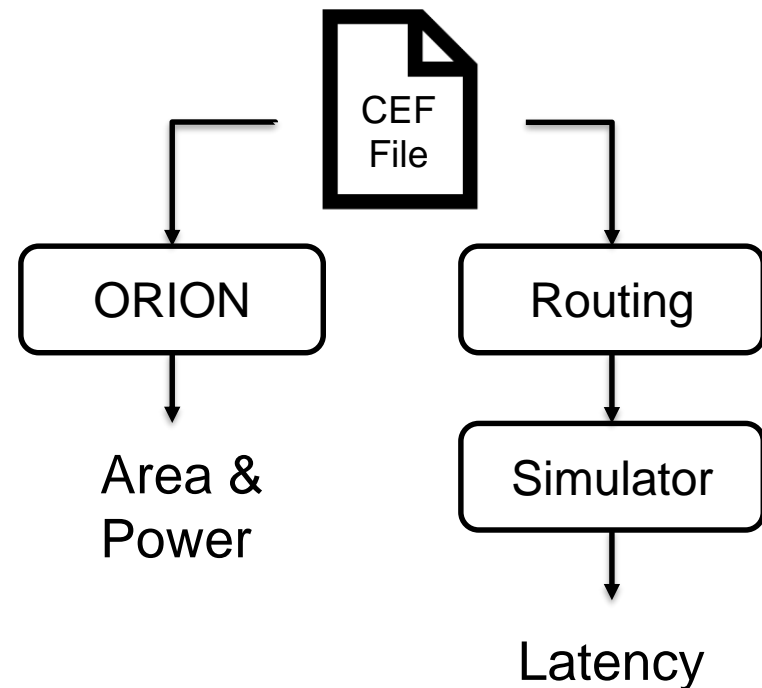
- Area and power model for NoCs

Routing:

- Shortest path routing algorithm

Simulator:

- Cycle-accurate SystemC simulation



Optimization of NoC topologies

Design task: Find the best-suited NoC topology for the application at hand

Objective:

$$\operatorname{argmax}_x Q(x, t)$$

Optimization methods:

- Gradient-based optimization **not applicable** as $Q(x, t)$ is not continuous
- Possible Heuristics:
 - Simulated Annealing
 - Genetic Algorithms
 - Monte-Carlo Tree Search

Recap: Markov Decision Process

Process description model for **sequential decision problems** in a **fully observable environment**

- **Sequential decision problem**

The utility of an action does not depend on a single decision but on the whole sequence of decisions

- **Fully observable environment**

The state of a system is known at all times

- **Definition***

S : set of possible states

A : set of possible actions

$t(s, a)$: state transition function returning state s' for the application of a on s

$Q(s)$: reward function

*Note: Definition slightly diverges from the Reinforcement Learning lecture

Monte Carlo Tree Search

- Tree search algorithm for Markov decision processes (MDP)
- Incrementally builds a search tree guided by previous explorations

Node

- State $s \in S$
- Quality of the state $Q(s)$
- Visit count of the state n_s

Edge

- Action $a \in A$

Objective

- Find action a^* in a state s which maximizes the future reward

Monte Carlo Tree Search and NoCs

Transfer NoC optimization problem into MDP

- **States S**
The complete design space of all NoC Topologies
- **Actions $A(s)$**
All valid modifications for a given NoC architecture s
- **Transition function $t(s, a)$**
Defines how an action modifies a NoC architecture
Formally expressed by graph rewriting
- **Reward function $Q(s)$**
Expresses the quality of a NoC architecture s

$$Q(s) = -1 * [w_l l_{avg}(s, t) + w_p P(s) + w_A A(s)]$$

DSE of NoC Architectures based on MCTS

Search Tree:

- Node: NoC Architecture s
- Edge: NoC modification a

Process:

Selection:


Select a node according to UCT function

$$UCT = Q_{ws}(s) + C_p \sqrt{\frac{2 \ln N(s_{root})}{N(s)}}$$

Exploitation

Exploration

Exploration
parameter



DSE of NoC Architectures based on MCTS

Search Tree:


- Node: NoC Architecture s
- Edge: NoC modification a

Process:

Expansion:

Apply n modifications on the selected node:

- Add/remove link
- Switch PEs
- Shift PE
- Merge routers



DSE of NoC Architectures based on MCTS

Search Tree:


- Node: NoC Architecture s
- Edge: NoC modification a

Process:

Simulation:

Evaluate the new NoC Architectures:

- Power and Area:
ORION
- Average Latencies:
Routing algorithm
SystemC simulation



DSE of NoC Architectures based on MCTS


Search Tree:

- Node: NoC Architecture s
- Edge: NoC modification a

Process:

Backpropagation:

Update the visit count of the ancestor architectures



DSE of NoC Architectures based on MCTS

Search Tree:

- Node: NoC Architecture s
- Edge: NoC modification a


Process:

Repeat until termination criteria is fulfilled

Time budget or simulation limit

Update root node

Select child with the best successor



DSE of NoC Architectures based on MCTS

Search Tree:

- Node: NoC Architecture s
- Edge: NoC modification a

Process:

Repeat till termination criteria is fulfilled

Update root node:


Select child with the best successor

Continue with the exploration

Expansion

Simulation

Backpropagation




DSE of NoC Architectures based on MCTS

Conclusion:

- ✓ **No domain knowledge required**
- ✓ **Traceability of the modifications**
 - Designers are able to trace the modifications and validate the result
 - Increases trust in the optimization tool
- ✗ **Low convergence speed**
 - Convergence to an optimized design is not given within reasonable time

DSE of NoC Architectures based on MCTS


- Analysis of the execution time identifies SystemC simulation as bottleneck
- **Objective:**
Estimate the SystemC simulation results
→ Neural Network



Outline

- Background
- Application of Neural Networks
 - Data Representation
 - Exploration of further neural network architectures
- Introduction to the ML framework

Data Representation Problem




Fully-connected neural network requires a vector as input data.

Hands-on exercise:

For the chosen problem, define the input vector of the neural network.


Mathematical Representation of NoCs

NoC Architecture




Adjacency Matrix

$$\begin{matrix} & PE_1 & PE_2 & PE_3 & R_1 & R_2 \\ \begin{matrix} PE_1 \\ PE_2 \\ PE_3 \\ R_1 \\ R_2 \end{matrix} & \begin{bmatrix} 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 \\ 1 & 0 & 0 & 0 & 1 \\ 0 & 1 & 1 & 1 & 0 \end{bmatrix}
 \end{matrix}$$




$$\begin{matrix} & PE_1 & PE_2 & PE_3 \\ \begin{matrix} PE_1 \\ PE_2 \\ PE_3 \end{matrix} & \begin{bmatrix} 0 & 40 & 100 \\ 20 & 0 & 0 \\ 0 & 10 & 0 \end{bmatrix} & \text{Mbit/s}
 \end{matrix}$$

Data Representation Problem




What if the network size changes?

Data Representation Problem



Data Representation Problem

NoC Architecture



Adjacency Matrix

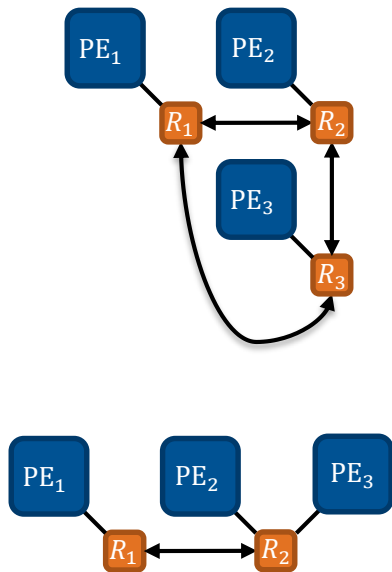
	PE_1	PE_2	PE_3	R_1	R_2	R_3
PE_1	0	0	0	1	0	0
PE_2	0	0	0	0	1	0
PE_3	0	0	0	0	0	1
R_1	1	0	0	0	1	1
R_2	0	1	0	1	0	1
R_3	0	0	1	1	1	0

Matrix becomes larger while input vector size is fixed.

How can we address this problem?

Zero Padding

NoC Architecture



Adjacency Matrix

$$\begin{array}{c}
 PE_1 \\
 PE_2 \\
 PE_3 \\
 R_1 \\
 R_2 \\
 R_3
 \end{array}
 \begin{bmatrix}
 PE_1 & PE_2 & PE_3 & R_1 & R_2 & R_3 \\
 0 & 0 & 0 & 1 & 0 & 0 \\
 0 & 0 & 0 & 0 & 1 & 0 \\
 0 & 0 & 0 & 0 & 0 & 1 \\
 1 & 0 & 0 & 0 & 1 & 1 \\
 0 & 1 & 0 & 1 & 0 & 1 \\
 0 & 0 & 1 & 1 & 1 & 0
 \end{bmatrix}$$

$$\begin{array}{c}
 PE_1 \\
 PE_2 \\
 PE_3 \\
 R_1 \\
 R_2 \\
 R_3
 \end{array}
 \begin{bmatrix}
 PE_1 & PE_2 & PE_3 & R_1 & R_2 & R_3 \\
 0 & 0 & 0 & 1 & 0 & 0 \\
 0 & 0 & 0 & 0 & 1 & 0 \\
 0 & 0 & 0 & 0 & 1 & 0 \\
 1 & 0 & 0 & 0 & 1 & 0 \\
 0 & 1 & 1 & 1 & 0 & 0 \\
 0 & 0 & 0 & 0 & 0 & 0
 \end{bmatrix}$$

- Chose a matrix size which fits the maximal supported number of routers and Pes
- Set the rows and columns of non-present routers/PEs to 0
→ matrices might become very sparse for large

Conclusion

- ✗ The number of features highly depends on the NoC size
 - Extensive zero padding would be required
 - ✗ Large NoCs require very big input vector ($> 10^3$)
 - High number of trainable parameters ($> 10^6$)
 - ✗ Sparse input vectors are difficult to train
 - ✗ Spatial invariance of the features is not considered
- Not suited for the problem at hand

Outline

- Background
- Application of Neural Networks
 - Data Representation
 - Exploration of further neural network architectures
- Introduction to the ML framework

Neural Network Architecture Exploration

Convolutional Neural Network

- ✓ Processes patterns in the input vector
- ✗ Designed for data in the Euclidean domain (i.e. not for graphs)
- ✗ The number of features highly depends on the NoC size
 - Extensive zero padding would be required
- ✗ Locality between neighboring routers vanishes after transformation into matrices
- Not suited for the problem at hand

Neural Network Architecture Exploration

Fully-Convolutional Neural Network

- ✓ Processes patterns in the input vector
- ✓ Processes arbitrary input vector sizes
- ✗ Designed for data in the Euclidean domain (i.e. not for graphs)
- ✗ Locality between neighboring routers vanishes after transformation into matrices
- Not suited for the problem at hand

Neural Network Architecture Exploration

Recurrent Neural Networks

- ✗ Designed for sequences of data
- Not suited for the problem at hand

Neural Network Architecture Exploration

Geometric Deep Learning[3]

- summarizes a set of neural network architectures for manifolds and graphs
- Still open field of research
- ✓ Designed for non-Euclidean domains i.e. graphs
- ✗ Graphs of arbitrary sizes are not supported
- Not suited for the problem at hand

[3] Federico Monti, “Geometric Deep Learning,” *Geometric Deep Learning*.
[Online]. Available: <http://geometricdeeplearning.com/>. [Accessed: 18-Jan-2019].

Outline

- Background
- Application of Neural Networks
 - Data Representation
 - Neural Network Architecture Exploration
 - Rewrite Problem
- Introduction to the ML framework

Rewrite Optimization Problem

Objective:

Estimate the average latency in a NoC

$$l_{avg} = \frac{1}{|P|} \sum_{p \in P} T_{rx}(p) - T_{tx}(p)$$


Rewritten Objective:

Estimate the average latency of each flow $f \in F$ individually

$$l_{avg} = \frac{1}{|F|} \sum_F \frac{bw_f}{BW_{NoC}} l_f$$

$$BW_{NoC} = \sum_F bw_f$$

Flows with higher BW send more packets and thus must be given a stronger weight



Design of the Recurrent Neural Network

Output:

The average latency of a flow f^*

Input:

Sequence of vectors describing the contentions along the routing path of flow f^*


$$S_{f^*} = (\underline{c}_0, \underline{c}_1, \dots, \underline{c}_n, \dots, \underline{c}_{N-1}); N = \text{number of hops}$$

Each vector \underline{c}_n collects the link utilization of all conflicting flows $F_{c_{f^*}}$ at the n^{th} hop of f^* :


$$\underline{c}_n = (u_{f^*}, u_{f_0}, u_{f_1}, \dots, u_{f_M}, 0, \dots, 0)^T; f_m \in F_{c_{f^*}} \text{ and } |\underline{c}_n| = 10$$

$$u_f := \frac{BW_f}{BW_{max}}$$

upper bound of
conflicting flows




Design of the Recurrent Neural Network




$$\begin{array}{c}
 PE_0 \\
 PE_1 \\
 PE_2 \\
 PE_3 \\
 PE_4
 \end{array}
 \begin{bmatrix}
 PE_0 & PE_1 & PE_2 & PE_3 & PE_4 \\
 0 & 0 & 0 & bw_{f_0} & 0 \\
 0 & 0 & 0 & bw_{f_1} & 0 \\
 0 & 0 & 0 & bw_{f_2} & 0 \\
 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & bw_{f_3} & 0
 \end{bmatrix}$$

Design of the Recurrent Neural Network



$$\begin{array}{c}
 PE_0 \\
 PE_1 \\
 PE_2 \\
 PE_3 \\
 PE_4
 \end{array}
 \begin{bmatrix}
 PE_0 & PE_1 & PE_2 & PE_3 & PE_4 \\
 0 & 0 & 0 & bw_{f_0} & 0 \\
 0 & 0 & 0 & bw_{f_1} & 0 \\
 0 & 0 & 0 & bw_{f_2} & 0 \\
 0 & 0 & 0 & 0 & 0 \\
 0 & 0 & 0 & bw_{f_3} & 0
 \end{bmatrix}$$



Recurrent Neural Network

Outline

- Background
- Application of Neural Networks
- Introduction to the Machine Learning Framework

Python IDE - Pycharm

Start Pycharm

```
cd /usr/local/labs/ML/  
./bin/ml pycharm &
```

Open Project

File → Open...

Change directory to `/usr/local/labs/ML/ab12cde/ml_eda`

Ok

Run or Debug Project

Run → Run.../Debug... → TaskX

Open a Terminal in Pycharm

View → Tool Windows → Terminal or Alt+F12

Open the Python Console in Pycharm

View → Tool Windows → Python Console

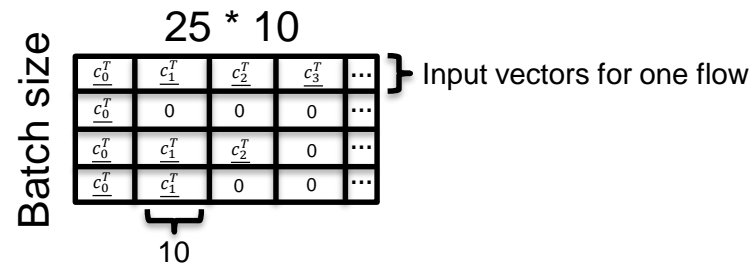
Project Structure

ml_eda

- data
 - features.npy
 - labels.npy

Training Data

- Features
 - 2d-array (batch size x concatenated input vectors)
 - Max number of contentions: 10
 - Max number of hops: 25



- Labels
 - 1d-array (batch size)
 - Average latency in cycles (from SystemC)

Note: First half of the data is for lowly utilized NoCs while the other half is for highly utilized NoCs!

Project Structure

ml_eda

- data
 - features.npy
 - labels.npy
- logs

Log-directory

- Output directory for all log-files generated during the training process including
 - Hyperparameters
 - Neural Network weights
 - Event files for Tensorboard

Project Structure

ml_edu

- data
 - features.npy
 - labels.npy
- logs
- **models**
 - genericRNN.py
 - simpleLSTM.py

Neural Network Models

- SimpleLSTM.py
 - Simple RNN model using LSTM cells
 - Used in Task 1
- GenericRNN.py
 - Parameterizable RNN
 - Used for the hyperparameter optimization in Task 2

Project Structure

ml_eda

- data
 - features.npy
 - labels.npy
- logs
- models
 - genericRNN.py
 - simpleLSTM.py
- `venv`

Virtual Python Environment

- Isolated python environment
- Stores ML libraries (e.g. keras and tensorflow)
- Do not modify this directory!

Project Structure

ml_eda

- data
 - features.npy
 - labels.npy
- logs
- models
 - genericRNN.py
 - simpleLSTM.py
- venv
- [HyperparameterOptimization.py](#)

Hyperparameter Optimization Class

- Wrapper Class for the scikit-optimize library
- Utilized in Task 2

Project Structure

ml_eda

- data
 - features.npy
 - labels.npy
- logs
- models
 - genericRNN.py
 - simpleLSTM.py
- venv
- HyperparameterOptimization.py
- [NeuralNetwork.py](#)

Neural Network Base Class

- Abstract base class for all models in *models*
- Main Methods:
 - *fitness(cls, hyperparameters)*
Trains a neural network for 5 epochs and returns the mean squared validation error
 - *split_data_set(cls, features, labels)*
Splits data set into training, validation and test data

Project Structure

ml_eda

- data
 - features.npy
 - labels.npy
- logs
- models
 - genericRNN.py
 - simpleLSTM.py
- venv
- HyperparameterOptimization.py
- NeuralNetwork.py
- [Task1.py](#)
- [Task2.py](#)

Task1

- Main class for the first task

Task2

- Main class for the second task

Recurrent Neural Networks in Keras

LSTM & GRU Class

- Implementation not suited for GPUs
- Important parameters
 - Units
 - Return_sequence

Masking Class

- Supports masking for input data of variable number of time steps
- Important parameters:
 - Mask_value
 - Input_shape

CuDNNLSTM & CuDNNGRU Class

- GPU only implementation of LSTM and GRU
- Do not support masking layer

Further Information: <https://keras.io/layers/recurrent/>

Task 1

a) Data pre-processing

- Load data set from the files „data/features.npy“ and „data/labels.npy“
- Separates the data set into 70% training, 15% validation and 15% test set using the method `split_data_set` in `NeuralNetwork.py`

b) Train the RNN

- Implement a LSTM network in the `models/SimpleRNN.py` class
- Tune the hyperparameters to obtain a mean absolute percentage error below 5%

Task 2

a) Hyperparameter Optimization

- Define a set of hyperparameters for your RNN.
- Implement a parameterizable RNN in GenericRNN.py.

b) Random Search

- Run a random search optimization for 15 iterations.
- Plot the results using Tensorboard and find the best network.

c) Bayesian Optimization

- Run a Bayesian optimization for 15 iterations (modify log directory).
- Plot the results using Tensorboard and find the best network.